

Search & AI


Stephen J. Guy

Many images from *Lavalle, Planning Algorithms* and *Peter Abbeel*

1

Search


- Very common in AI
 - “AI is search”
 - Game trees, decision trees, puzzle solving
- We’ll see the same things in path planning



2

2

Search Trees





- “What if” analysis of plans and outcomes
- **Root:** Initial state
- **Children:** Successors following an action
- **Node:** Contain states; node's parents correspond to plans reaching that state
- **Fringe:** Nodes with unexplored action

3

3

Depth First Search

- Strategy:
 - Expand deepest node
- Implementation:
 - Fringe is LIFO stack





Q: What can go wrong?

4

4

Breadth First Search

- Strategy:
 - Expand shallowest node
- Implementation:
 - Fringe is FIFO queue



5

5


Breadth-first vs Depth-first

- DFS:
 - Less Memory
 - Will not find optimal solution
 - Gets stuck in loops and infinite paths
- Breadth First:
 - Lots of memory required
 - Finds solution with least nodes
 - Handles loops & infinite paths
- Hybrid Options...

6

6

Costs on Actions





- What if different actions have differing costs
- BFS finds least num. of actions, not least cost path

7

7

Uniform Cost Search

- Strategy:
 - Expand node w/ cheapest total
- Implementation:
 - Fringe is a priority queue



8

8

Aside: Priority Queue

- Priority queue allows insert of (key, value) pairs

<code>pq.push(key, value)</code>	inserts (key, value) into the queue.
<code>pq.pop()</code>	returns the key with the lowest value, and removes it from the queue.


- Insertions normally $O(\log n)$
- Decrease value: `pop()`, `push(key, val-1)`
- Priority Queues used in cost-sensitive search methods

9

9

Uniform Cost Search: Analysis

- UCS: Explore increasing cost contours
- Advantages:
 - UCS is complete and optimal
- Issue:
 - Explores in every "direction" equally
 - Uses no info. about goal




10

10

Heuristics


- Straight-line heuristic



11

UCS + Heuristic


- Uniform-cost orders by path-cost so far
 - Cost-to-node: $g(n)$
- A* Search orders by total estimated cost $f(n)$
 - $f(n) = g(n) + h(n)$



12

UCS vs A* Contours

- Uniform cost expands in all directions
- A* expands mainly towards the goal
 - Hedges by expanding outward as needed




13

13

Terminating A*


- Can we stop when we push the goal onto the PQ?
- No – Terminate on dequeuing a goal



14

14

Is A* Optimal?



- What went wrong?
 - Actual cost to goal < estimate cost to goal
- Heuristic must be less than the actual cost!
 - A* assumes it will take at least $h()$ to get to goal

15

15



Admissible Heuristic

- A heuristic h is **admissible** if

$$h(n) \leq h^*(n)$$

where $h^*(n)$ is the true cost to the nearest goal

- Examples:



- Finding admissible heuristic is main part of using A* in practice

16

16

Optimality

- Terminating A*
 - A* is done when the goal is the lowest thing on the priority queue.
- A* finds optimal path if heuristic is admissible!

17

17

Classic Example

- 8 Puzzle
- Empty space can slide up, down, left, or right
- Ideas for heuristics?
 - Num displace tiles
 - Total Manhattan dist.

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8


Goal State

Average nodes expanded when optimal path has length...			
	...4 steps	...8 steps	...12 steps
UCS	112	6,300	3.6×10^6
TILES	13	39	227
MANHATTAN	12	25	73

18

Comparison

- UCS
- A*
- Greedy
 - *Inadmissible*



19

Weighed A* $f = g + \epsilon h$


- Weighted A*
 - Expands nodes in order of $f(n) = g(n) + \epsilon h(n)$
 - $\epsilon > 1$
 - Bias towards goal nodes

20

20

Weighted A*: $\epsilon = 0$


- $\epsilon = 0 \rightarrow$ Uniform cost search



21

Weighted A*: $\epsilon = 1$


- $\epsilon = 1 \rightarrow A^*$



22

Weighted A*: $\epsilon > 1$

- $\epsilon > 1 \rightarrow$ Targeted path towards goal



23

23

Anytime Algorithms

24

24

Weighted A* $f = g + \epsilon h : \epsilon > 1$

- Trades off optimality for speed
- ϵ -suboptimal:
 - $\text{cost}(\text{solution}) \leq \epsilon * \text{cost}(\text{optimal solution})$
 - Test your understanding by trying to prove this!
- In many domains, it has been shown to be orders of magnitude faster than A* !
- Research challenge is to find heuristics with shallow local minima

25

25

Anytime Algorithm


- Algorithm that returns valid answer when interrupted mid run
- Use in robotics
 - Robot can make best current decision as it moves given current information
- Use in games/VR
 - NPCs can make the best decision given time budget for AI

26

26

Anytime A*

- Weighted A*
 - Trades off speed for optimality
 - ϵ -suboptimal
- Anytime A*
 - Run weighted A* large ϵ
 - Rerun with decreasing ϵ to refine path



27

A* Variants

- **Anytime Repairing A* (ARA*)** [Likhachev, Gordon, Thrun 2004]
 - Efficient version of Anytime A* that reuses state values within each iteration
- **Anytime Nonparametric A* (ANA*)** [van den Berg, Shah, Huang, Goldberg, 2011]
 - Provides the “right” next ϵ to ARA*
- **Lifelong Planning A* (LPA*)**
 - Modified A* to allow updates in weights (e.g., passage becomes blocked)

28

28

D*: A*, but Goal to Start




- A* with consistent heuristic finds shortest path from goal to start
- Flip start and goal
 - Shortest path from any point to start
- Policy for any point in space
 - Can account for noise/errors in motion

29

29

D* Lite – Online Planning

- Combine
 - Lifelong Planning A*
 - Flip Start and Goal
 - Some other optimization
- Result
 - Can account for new observations as updated weights on edges
 - e.g. blocked paths



30

30

D* in practice

- **Field D*** implemented on Mars rovers Spirit and Opportunity
- **Anytime D*** used in DARA urban challenge winner



31

Research/Project Ideas

- Can you use LPA*/D* ideas to create NPCs which explore games worlds intelligently
 - Allow dynamic environments
 - NPCs in "unknown" environments
 - Play hide-and-seek



32

32

Terms

- Search Tree (Root, Node, Child, Fringe)
- Depth First Search (DFS)
- Breadth First Search (BFS)
- Heuristics
- A*
- Admissibility
- Anytime Algorithms
- Policies

33

33